



convolutional neural network

Hand Gesture Classification



Project Goals

- ☐ Create own CNN model
- ☐ Train network capable of classifying hand gestures with sufficient accuracy
- ☐ Provide analysis and testing for the model

Tech

- ☐ PyTorch 2.0
- ☐ Python 3.10



Dataset - HaGRID

Original dataset contains 552,992
FullHD RGB images divided into
18 classes of gestures




- ❑ For the needs of the program a subset of the dataset is used.
- ❑ Subset consist of cropped images of hands only
- ❑ Number of images is reduced to ~10k

<https://github.com/hukenovs/hagrid>



Data preparation

- ❑ Consistent size: 64-256px
- ❑ Pixel normalization
- ❑ Data is splitted onto test(10%) and train data (90%).
- ❑ Classes divided in folders on a hard drive
- ❑ Data is shuffled, the seed of the shuffle is const.



Two sides of the approach

Two different architectures

TinyVGG

Compact version of VGG, easy to implement, only few layers.

- Straightforward and quick,
- Clear,
- Lightweight,
- Ability to learn NN features,

RasNet18

Powerful deep network, consists of 18 layers. One of RasNet variants.

- Employs residual connections,
- Fully connected layers,
- High accuracy
- Long training time

Usage of two architectures allows for better understanding and greater comparison capabilities



Model 1 - TinyVGG

- ❑ Based on VGG architecture with own tweaks
- ❑ On this example:
 - 4 convolution layers
 - 2 pooling layers
 - no extra fully connected layers
 - Uses ReLU for activation

```
##### MODEL #####
class CustomTinyVGG(nn.Module):
    def __init__(self, input_shape: int, hidden_units: int, output_shape: int) -> None:
        super().__init__()
        self.convBlock_1 = nn.Sequential(
            nn.Conv2d(in_channels=input_shape,
                      out_channels=hidden_units,
                      kernel_size=6,
                      stride=1,
                      padding=2),
            nn.ReLU(),
            nn.Conv2d(in_channels=hidden_units,
                      out_channels=hidden_units,
                      kernel_size=5,
                      stride=1,
                      padding=1),
            nn.ReLU(),
            nn.MaxPool2d(kernel_size=2,
                         stride=2)
        )
        self.convBlock_2 = nn.Sequential(
            nn.Conv2d(hidden_units, hidden_units, kernel_size=4, padding=1),
            nn.ReLU(),
            nn.Conv2d(hidden_units, hidden_units, kernel_size=4, padding=1),
            nn.ReLU(),
            nn.MaxPool2d(2)
        )
        self.classifier = nn.Sequential(
            nn.Flatten(),
            nn.Linear(in_features=30*30*hidden_units, out_features=output_shape)
        )

    def forward(self, mod: torch.Tensor):
        mod = self.convBlock_1(mod)
        mod = self.convBlock_2(mod)
        mod = self.classifier(mod)
        return mod
```

```
woq = 2671*CT9227476L(woq)
```

Model 2 – RasNet18

layer name	output size	18-layer	34-layer	50-layer	101-layer	152-layer
conv1	112×112	7×7, 64, stride 2				
conv2_x	56×56	3×3 max pool, stride 2				
		$\begin{bmatrix} 3 \times 3, 64 \\ 3 \times 3, 64 \end{bmatrix} \times 2$	$\begin{bmatrix} 3 \times 3, 64 \\ 3 \times 3, 64 \end{bmatrix} \times 3$	$\begin{bmatrix} 1 \times 1, 64 \\ 3 \times 3, 64 \\ 1 \times 1, 256 \end{bmatrix} \times 3$	$\begin{bmatrix} 1 \times 1, 64 \\ 3 \times 3, 64 \\ 1 \times 1, 256 \end{bmatrix} \times 3$	$\begin{bmatrix} 1 \times 1, 64 \\ 3 \times 3, 64 \\ 1 \times 1, 256 \end{bmatrix} \times 3$
conv3_x	28×28	$\begin{bmatrix} 3 \times 3, 128 \\ 3 \times 3, 128 \end{bmatrix} \times 2$	$\begin{bmatrix} 3 \times 3, 128 \\ 3 \times 3, 128 \end{bmatrix} \times 4$	$\begin{bmatrix} 1 \times 1, 128 \\ 3 \times 3, 128 \\ 1 \times 1, 512 \end{bmatrix} \times 4$	$\begin{bmatrix} 1 \times 1, 128 \\ 3 \times 3, 128 \\ 1 \times 1, 512 \end{bmatrix} \times 4$	$\begin{bmatrix} 1 \times 1, 128 \\ 3 \times 3, 128 \\ 1 \times 1, 512 \end{bmatrix} \times 8$
conv4_x	14×14	$\begin{bmatrix} 3 \times 3, 256 \\ 3 \times 3, 256 \end{bmatrix} \times 2$	$\begin{bmatrix} 3 \times 3, 256 \\ 3 \times 3, 256 \end{bmatrix} \times 6$	$\begin{bmatrix} 1 \times 1, 256 \\ 3 \times 3, 256 \\ 1 \times 1, 1024 \end{bmatrix} \times 6$	$\begin{bmatrix} 1 \times 1, 256 \\ 3 \times 3, 256 \\ 1 \times 1, 1024 \end{bmatrix} \times 23$	$\begin{bmatrix} 1 \times 1, 256 \\ 3 \times 3, 256 \\ 1 \times 1, 1024 \end{bmatrix} \times 36$
conv5_x	7×7	$\begin{bmatrix} 3 \times 3, 512 \\ 3 \times 3, 512 \end{bmatrix} \times 2$	$\begin{bmatrix} 3 \times 3, 512 \\ 3 \times 3, 512 \end{bmatrix} \times 3$	$\begin{bmatrix} 1 \times 1, 512 \\ 3 \times 3, 512 \\ 1 \times 1, 2048 \end{bmatrix} \times 3$	$\begin{bmatrix} 1 \times 1, 512 \\ 3 \times 3, 512 \\ 1 \times 1, 2048 \end{bmatrix} \times 3$	$\begin{bmatrix} 1 \times 1, 512 \\ 3 \times 3, 512 \\ 1 \times 1, 2048 \end{bmatrix} \times 3$
	1×1	average pool, 1000-d fc, softmax				
FLOPs		1.8×10^9	3.6×10^9	3.8×10^9	7.6×10^9	11.3×10^9

- ❑ RasNet variant with 18 layers in total
- ❑ Fully connected layers
- ❑ Also uses ReLU for activation
- ❑ Normalization to stabilize training process using BatchNorm2d

ures for ImageNet. Building blocks are shown in brackets (see also Fig. 5), with the numbers of block

Source: https://pytorch.org/hub/pytorch_vision_resnet/



Training Process

Parameters that were adjusted to get best effect:

- Learning rate
- Optimizer algorithm
- Batch size & Epochs
- Number of Convolutional layers
- Padding and Kernels sizes
- Dropouts
- Number of Pooling layers
- Activation functions
- Weight Decay
- Image dimensions
- Overall model complexity

Consts





- ☐ Seed
- ☐ Dataset Split
- ☐ Device (GPU)
- ☐ Image channels
- ☐ Classes
- ☐ Loss function

Trained on up to ~10000 images



TinyVGG

Current Results





~75%
accuracy

~80%
accuracy



RasNet18

Current Results




~90%
accuracy

Pred: fist | Prob: 0.986



Pred: two_up_inverted | Prob: 0.586



Pred: peace_inverted | Prob: 0.570



Pred: palm | Prob: 0.572



Pred: stop_inverted | Prob: 0.898



Pred: call | Prob: 0.997




Pred: dislike | Prob: 1.000



Pred: rock | Prob: 0.644



Pred: three2 | Prob: 0.998



Examples





Not so accurate guesses

Pred: three | Prob: 0.817



Pred: four | Prob: 0.289



Pred: dislike | Prob: 0.465



Pred: one | Prob: 0.701



<< Faulty data



The End

Some hands-related proverbs

Two watermelons can't be held in one hand. ~ Afghan Proverb

A gentle hand may lead even an elephant by a single hair. ~ Iranian Proverb

You can't use your hand to force the sun to set. ~ Nigerian Proverb

Promises have legs. Only a gift has hands. ~ German Proverb

by: Adrian Zaręba